

Table of Contents

- [Base Notation](#)
- [Base Terminology](#)
- [Environments](#)
- [Agents](#)

Base Notation

k : number of bandit arms, or size of the bandit environment

t : discrete time step

a : denotes an action. $a \in k$

A_t : denotes the action taken at time point t

$R_{t,a}$: realized (ex-post) reward at time point t , given that $A_t = a$

$Q_t(a)$: reward estimate for action a at time point t

$q_*(a) = E(R_t | A_t = a)$: expected reward of action a at time point t

$N_t(a)$: number of times action a has been selected until time point t

Base Terminology

Agent: The instance/robot to solve the problem.

Environment: The problem to solve.

Exploitation (greedy action): Choosing the action with the highest estimated value $\operatorname{argmax}_a Q_t(a)$, for the purpose of maximising the (expected) total reward.

Exploration: Choosing non-greedy actions for optimising/updating the estimates $Q_t(a)$ for each a which is *not* $\operatorname{argmax}_a Q_t(a)$. Exploration has the purpose of calculating better estimates for non-greedy actions, with the aim to minimize $|q_*(a) - Q_t(a)|$.

Valuation methods: Valuation methods define how the reward estimates $Q_t(a)$ are calculated.

valuation_method = average: Considering a single action: $Q_{n+1} = \frac{R_1 + R_2 + \dots + R_n}{n}$, where n denotes the n -th selection of the respective action. This can be re-written as

$$Q_{n+1} = Q_n + \frac{1}{n} [R_n - Q_n]$$

valuation_method = weighted: Considering a single action: $Q_{n+1} = Q_n + \alpha [R_n - Q_n]$

Environments

At the moment, the repository contains 3 bandit environments.

1. Stochastic Bandit Environment `stochastic_bandit`

The environment `stochastic_bandit` is a stationary bandit problem, where the reward of each arm is a random variable coming from either a gaussian normal or a log-normal distribution.

When initialising the environment, the expected rewards for each action $q_{(a)}$ * from a normal distribution, where μ and σ are argument inputs of the environment class (`stochastic_moments = [0, 1]` per default.)

Arguments

`size`: the number of bandit arms k

`reward_distribution`: defines the theoretical distribution of rewards. Options: `["gaussian", "lognormal"]`.

`stochastic_moments`: Defines the first and second moment of the normal

distribution the expected rewards $q_*(1), \dots, q_*(k)$ are sampled from ($\mu = 0$ and $\sigma = 1$ by default).

Explanation, 5-armed bandit case: If the environment is initialised with `stochastic_moments = [1, 2]`, the expected returns for each of the 5 reward distributions $q_*(1), \dots, q_*(5)$ are random variables from a normal distribution with $\mu = 1$ and $\sigma = 2$ ($N(1, 2)$). After the initialisation, let $q_*(5) = 0.2$. Subsequently, $R_{t,5} \sim N(0.2, 1)$ in case of normally distributed rewards ($\sim \log N(0.2, 1)$ applies if the user chose the log-normal distribution).

2. Custom Stochastic Bandit Environment `stochastic_bandit_custom`

The environment `stochastic_bandit_custom` is a stationary bandit problem. This environment grants more flexibility for defining the underlying problem, compared to `stochastic_bandit`.

Through the input argument `mu` and `sigma`, the user can customise the theoretical distributions which are used for reward sampling.

For `stochastic_bandit_custom`, the reward distribution are completely customisable, while `stochastic_bandit` samples the parameters of the rewards distributions during initialisation (i.e., the distribution parameters are stochastic itself).

Arguments

`size`: the number of bandit arms k

`mu`: an array of central values for the reward distributions, which are similar to the expected returns for each action $q_*(a)$.

`sigma`: an array of second stochastic moments for each reward distribution.

`reward_distribution`: defines the theoretical distribution of rewards. Options: ["gaussian", "lognormal"].

`stochastic_moments`: Defines the first and second moment of the normal distribution the expected rewards $q_*(1), \dots, q_*(k)$ are sampled from ($\mu = 0$ and $\sigma = 1$ by default).

Explanation, 2-armed bandit case: If the environment is initialised with `mu = [0, 0.5]` and `sigma = [2, 1]`, the rewards for both actions follow $R_{t,1} \sim N(0, 2)$ and $R_{t,2} \sim N(0.5, 1)$ if `reward_distribution = "gaussian"`.

3. Adversarial (non-stationary) Bandit `adversarial_bandit`

The environment `adversarial_bandit` is a non-stationary version of the `stochastic_bandit` environment.

Arguments

`size`: the number of bandit arms k

`reward_distribution`: defines the theoretical distribution of rewards. Options: ["gaussian", "lognormal"].

`stochastic_moments`: Defines the first and second moment of a normal distribution which is used for sampling the stationary component of the expected rewards $q_*(1), \dots, q_*(k)$ ($\mu = 0$ and $\sigma = 1$ by default).

Explanation, 3-armed bandit case: Given is an initiation with `stochastic_moments = [1, 2]`. The expected rewards are then defined as $E(R_t | A_t = a) \sim N(1, 2)$ for every action $a \in [1, \dots, k]$, implying that the expected reward $q_*(a)$ is constant over time and not dependent on t (note that $k=3$ for this example). Let $q_*(3) = 1.2$. We can write $R_{t,3} \sim N(1.2, 0.1t)$, indicating that rewards will fluctuate stronger with increasing t .

Agents

There are currently 3 different agents available in the repository. Generally, agents differ on how to select the next action a .

1. **Random agent RandomAgent**

The RandomAgent class selects action a randomly ($a \sim U(1, k)$).

Arguments

env: a bandit environment.

q_initialisation: an array of length k with initial q-estimates. None indicates initiation with zeros.

valuation_method: Defines how q-estimates are calculated. Must be either average or weighted.

alpha: Scale parameter, weights importance of last reward observation when calculating q-estimates. Required if valuation_method = weighted.

name: Custom name for your agent (optional).

2. **Epsilon-greedy agent EpsilonGreedyAgent**

This agent selects a greedy action with probability $1 - \epsilon$ and subsequently choose a (random) non-greedy action with probability ϵ . The probability for exploitation can be written as

$$P\left(A_t = \underset{a}{\operatorname{argmax}} Q_t(a)\right) = 1 - \epsilon$$

Arguments

env: a bandit environment.

epsilon: Constant, input for ϵ (see above)

q_initialisation: an array of length k with initial q-estimates. None indicates initiation with zeros.

valuation_method: Defines how q-estimates are calculated. Must be either average or weighted.

alpha: Scale parameter, weights importance of last reward observation when calculating q-estimates. Required if valuation_method = weighted.

name: Custom name for your agent (optional).

3. **Upper-Confidence-Bound Agent UCBAgent**

The UCB-Agent takes into account how “confident” it is in his Q-estimates by taking into account how often the respective action has been selected. The selection method can be written as:

$$A_t = \underset{a}{\operatorname{argmax}} \left[Q_t(a) + c \sqrt{\frac{\ln t}{N_t(a)}} \right]$$

, where $N_t(a)$ is the number of times action a has been selected until time point t .

Thus, the term $\sqrt{\frac{\ln t}{N_t(a)}}$ (the uncertainty term) puts the number of times an action a has been selected in relation to the total number of steps t , resulting in an addition to its estimate $Q_t(a)$ for actions which are less frequently selected. The constant c is a fixed parameters for scaling the uncertainty term.

Arguments

env: a bandit environment.

c: Constant, input for c (see above)

q_initialisation: an array of length k with initial q-estimates. None indicates initiation with zeros.

valuation_method: Defines how q-estimates are calculated. Must be either average or weighted.

`alpha`: Scale parameter, weights importance of last reward observation when calculating q-estimates. Required if `valuation_method = weighted`.
`name`: Custom name for your agent (optional).